

## Answer Sheet :

### Machine Learning Assignment Submission - (Parikshit Prajapati)

Name : Prajapati Parikshit Virabhai

Phone No. : 9998428706

Email : parikshitprajapati2001@gmail.com

Use the [Oil Spill Dataset](#) and solve the following question by using the dataset, to download the dataset click on the dataset name.

#### **About Dataset :**

The dataset was developed by starting with satellite images of the ocean, Some of which contain an oil spill and some that do not.

Images were split into sections and processed using computer vision algorithms to provide a vector of features to describe the contents of the image section or patch.

The task is, given a vector that describes the contents of a patch of a satellite

image, then predicts whether the patch contains an oil spill or not, e.g. from

the illegal or accidental dumping of oil in the ocean.

There are two classes and the goal is to distinguish between spill and non-spill using the features of a given ocean patch.

- **Non-Spill: negative case, or majority class.**
- **Oil Spill: positive case, or minority class.**

There are a total of 50 Columns in the Dataset, the output column is named as a target.

## QUESTIONS:

**Q1)** Download the Oil Spill Dataset and perform Data cleaning and Data Pre-Processing if Necessary.

**Q2)** Use various methods such as Handling null values, One-Hot Encoding, Imputation, and Scaling of Data Pre-Processing where necessary.

**Q3)** Derive some insights from the dataset.

**Q4)** Apply various Machine Learning techniques to predict the output in the target column, make use of Bagging and Ensemble as required, and find the best model by evaluating the model using Model evaluation techniques.

**Q5)** Save the best model and Load the model

**Q6)** Take the original data set and make another dataset by randomly picking 20 data points from the oil spill dataset and applying the save model to the same.

## INPUT AND OUTPUT :

**Q1) Download the Oil Spill Dataset and perform Data cleaning and Data:**

**Input :**

```
Import libraries :

In [59]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.metrics import accuracy_score, classification_report
import joblib

Load the Dataset

In [61]: df = pd.read_csv('oil_spill.csv')
df.head()

Out[61]:
```

	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_10	...	f_41	f_42	f_43	f_44	f_45	f_46	f_47	f_48	f_49	target
0	1	2558	1506.09	456.63	90	6395000	40.88	7.89	29780.0	0.19	...	2850.00	1000.00	763.16	135.46	3.73	0	33243.19	65.74	7.95	1
1	2	22325	79.11	841.03	180	55812500	51.11	1.21	61900.0	0.02	...	5750.00	11500.00	9593.48	1648.80	0.60	0	51572.04	65.73	6.26	0
2	3	115	1449.85	608.43	88	287500	40.42	7.34	3340.0	0.18	...	1400.00	250.00	150.00	45.13	9.33	1	31692.84	65.81	7.84	1
3	4	1201	1562.53	295.65	66	3002500	42.40	7.97	18030.0	0.19	...	6041.52	761.58	453.21	144.97	13.33	1	37696.21	65.67	8.07	1
4	5	312	950.27	440.86	37	780000	41.43	7.03	3350.0	0.17	...	1320.04	710.63	512.54	109.16	2.58	0	29038.17	65.66	7.35	0

5 rows × 50 columns

Data Cleaning and Pre-processing:

Data Cleaning and Pre-processing:

```
In [63]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 937 entries, 0 to 936
Data columns (total 50 columns):
#   column      Non-Null Count  Dtype
---  ---
0   f_1          937 non-null      int64
1   f_2          937 non-null      int64
2   f_3          937 non-null      float64
3   f_4          937 non-null      float64
4   f_5          937 non-null      int64
5   f_6          937 non-null      int64
6   f_7          937 non-null      float64
7   f_8          937 non-null      float64
8   f_9          937 non-null      float64
9   f_10         937 non-null      float64
10  f_11         937 non-null      float64
11  f_12         937 non-null      float64
12  f_13         937 non-null      float64
13  f_14         937 non-null      float64
14  f_15         937 non-null      float64
15  f_16         937 non-null      float64
16  f_17         937 non-null      float64
17  f_18         937 non-null      float64
18  f_19         937 non-null      float64
19  f_20         937 non-null      float64
20  f_21         937 non-null      float64
21  f_22         937 non-null      float64
22  f_23         937 non-null      int64
23  f_24         937 non-null      float64
24  f_25         937 non-null      float64
25  f_26         937 non-null      float64
26  f_27         937 non-null      float64
```

```
In [64]: df.isnull().sum()

Out[64]:
f_1      0
f_2      0
f_3      0
f_4      0
f_5      0
f_6      0
f_7      0
f_8      0
f_9      0
f_10     0
f_11     0
f_12     0
f_13     0
f_14     0
f_15     0
f_16     0
f_17     0
f_18     0
f_19     0
f_20     0
f_21     0
f_22     0
f_23     0
f_24     0
f_25     0
f_26     0
f_27     0
f_28     0
f_29     0
f_30     0
f_31     0
f_32     0
f_33     0
f_34     0
```

Describe Database:

```
In [65]: df.describe()

Out[65]:
```

	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_10	...	f_41
count	937.000000	937.000000	937.000000	937.000000	937.000000	9.370000e+02	937.000000	937.000000	937.000000	937.000000	...	937.000000
mean	81.588047	332.842049	698.707086	870.992209	84.121665	7.696964e+05	43.242721	9.127887	3940.712914	0.221003	...	933.928677
std	64.976730	1931.938570	599.965577	522.799325	45.361771	3.831151e+06	12.718404	3.588878	8167.427625	0.090316	...	1001.681331
min	1.000000	10.000000	1.920000	1.000000	0.000000	7.031200e+04	21.240000	0.830000	667.000000	0.020000	...	0.000000
25%	31.000000	20.000000	85.270000	444.200000	54.000000	1.250000e+05	33.650000	6.750000	1371.000000	0.160000	...	450.000000
50%	64.000000	65.000000	704.370000	761.280000	73.000000	1.863000e+05	39.970000	8.200000	2090.000000	0.200000	...	685.420000
75%	124.000000	132.000000	1223.480000	1260.370000	117.000000	3.304680e+05	52.420000	10.760000	3435.000000	0.260000	...	1053.420000
max	352.000000	32389.000000	1893.080000	2724.570000	180.000000	7.131500e+07	82.640000	24.690000	160740.000000	0.740000	...	11949.330000

8 rows × 50 columns

```
In [66]: df.duplicated().sum()

Out[66]: 0

In [70]: df.dtypes

Out[70]:
f_1      int64
f_2      int64
f_3      float64
f_4      float64
f_5      int64
f_6      int64
f_7      float64
```

## Q2) Use various methods such as Handling null values, One-Hot Encoding, Imputation, and Scaling of Data Pre-Processing where necessary.

### Feature Scaling:

```
In [71]: features = df.loc[:, 'f_2': 'f_49']
# features
scaler = StandardScaler()

scaled_features = scaler.fit_transform(features)
df.loc[:, 'f_2': 'f_49'] = scaled_features
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_12512\3951132622.py:6: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will attempt to set the values inplace instead of always setting a new array. To retain the old behavior, use either `df[df.columns[i]] = newvals` or, if columns are non-unique, `df.isetitem(i, newvals)`

```
df.loc[:, 'f_2': 'f_49'] = scaled_features
```

### Correlation Analysis:

```
In [72]: correlation_matrix = df.corr()
```

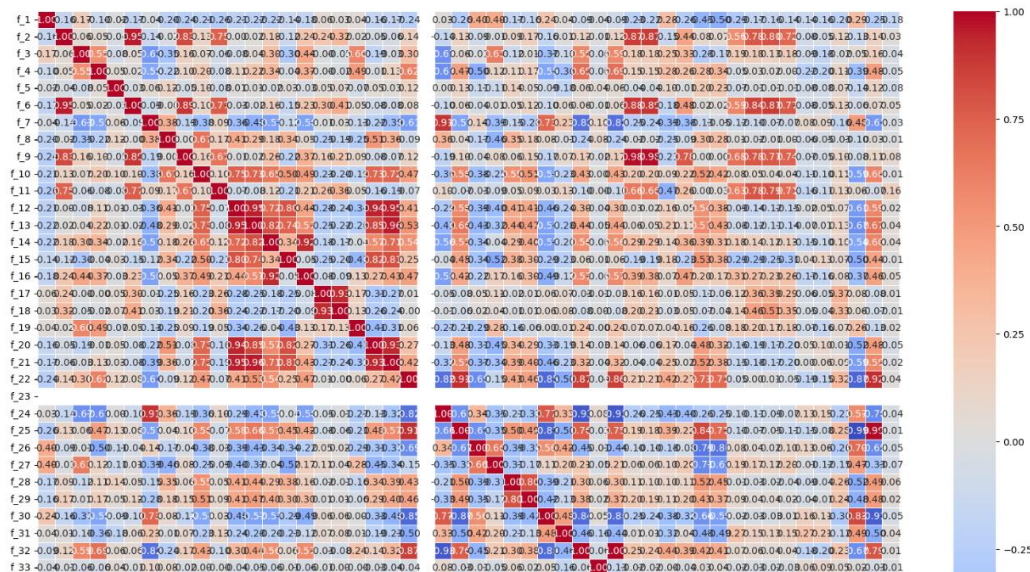
```
In [73]: correlation_matrix
```

```
Out[73]:
```

	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_10	...	f_41	f_42	f_43	f_44
f_1	1.000000	-0.155581	0.172017	-0.104116	-0.017025	-0.169533	-0.037412	-0.204983	-0.244551	-0.214447	...	-0.286190	-0.167466	-0.156916	-0.141792
f_2	-0.155581	1.000000	0.058390	0.052638	-0.036870	0.953947	-0.136761	-0.016822	0.829978	0.128465	...	0.555154	0.777807	0.800939	0.716496
f_3	0.172017	0.058390	1.000000	0.549510	-0.082764	0.050795	-0.627934	-0.349541	0.158686	0.073794	...	0.186920	0.178287	0.129653	0.176883
f_4	-0.104116	0.052638	0.549510	1.000000	0.048847	0.024693	-0.546205	-0.222063	0.097683	0.202167	...	-0.046934	0.032402	0.022234	0.000664
f_5	-0.017025	-0.036870	-0.082764	0.048847	1.000000	-0.028431	0.059128	0.123814	-0.047879	0.098573	...	-0.066930	-0.014877	-0.013742	-0.012346
f_6	-0.169533	0.953947	0.050795	0.024693	-0.028431	1.000000	-0.093589	-0.001395	0.894150	0.097449	...	0.594273	0.844597	0.868353	0.770044
f_7	-0.037412	-0.136761	-0.627934	-0.546205	0.059128	-0.093589	1.000000	0.381206	-0.188076	-0.380340	...	-0.115014	-0.100003	-0.074308	-0.073751
f_8	-0.204983	-0.016822	-0.349541	-0.222063	0.123814	-0.001395	0.381206	1.000000	0.001073	0.670628	...	0.013476	-0.015712	-0.013193	0.002439
f_9	-0.244551	0.829978	0.158686	0.097683	-0.047879	0.894150	-0.188076	0.001073	1.000000	0.164098	...	0.675610	0.784833	0.770129	0.736075
f_10	-0.214447	0.128465	0.073794	0.202167	0.098573	0.097449	-0.380340	0.670628	0.164098	1.000000	...	0.082449	0.052518	0.043116	0.042269
f_11	-0.261624	0.745590	-0.064076	-0.082742	-0.075843	0.765628	0.093376	0.167904	0.671358	0.102331	...	0.630674	0.782581	0.790649	0.710990
f_12	-0.209190	0.004035	-0.081738	0.106767	0.009470	-0.029363	-0.363593	0.406409	-0.008391	0.747509	...	-0.088211	-0.135129	-0.121701	-0.147694
f_13	-0.222342	0.020195	0.042723	0.224342	0.013574	-0.017706	-0.481003	0.289904	0.018342	0.730810	...	-0.084692	-0.120182	-0.109534	-0.140570
f_14	-0.2220721	0.176080	0.299324	0.335270	-0.016254	0.155767	-0.574566	0.178362	0.261617	0.652360	...	0.177034	0.141294	0.117372	0.130096
f_15	-0.137901	-0.118317	-0.301641	-0.039329	0.028305	-0.147712	-0.115334	0.335692	-0.215468	0.502049	...	-0.292963	-0.293204	-0.250771	-0.308273
f_16	-0.178220	0.235500	0.439603	0.372116	-0.029425	0.226015	-0.563544	0.051995	0.365164	0.487945	...	0.305778	0.269345	0.227190	0.262997
f_17	0.056430	0.237388	-0.003753	-0.000815	0.045836	0.302462	-0.008360	-0.245330	0.160027	-0.231361	...	0.119187	0.361130	0.392898	0.287938
f_18	0.027526	0.321276	-0.046857	-0.020119	0.065762	0.406917	0.027642	-0.188000	0.207135	-0.196430	...	0.144958	0.459463	0.509775	0.353361

## Q3) Derive some insights from the dataset.

```
In [74]: plt.figure(figsize = (20,16))
sns.heatmap(correlation_matrix, annot = True, cmap = 'coolwarm', fmt = '.2f', linewidths = .5)
plt.show()
```





#### Drop the column:


```
In [78]: df.drop(columns=['f_23'], inplace=True)
```

```
In [79]: df.columns
```


```
Out[79]: Index(['f_1', 'f_2', 'f_3', 'f_4', 'f_5', 'f_6', 'f_7', 'f_8', 'f_9', 'f_10',  
              'f_11', 'f_12', 'f_13', 'f_14', 'f_15', 'f_16', 'f_17', 'f_18', 'f_19',  
              'f_20', 'f_21', 'f_22', 'f_24', 'f_25', 'f_26', 'f_27', 'f_28', 'f_29',  
              'f_30', 'f_31', 'f_32', 'f_33', 'f_34', 'f_35', 'f_36', 'f_37', 'f_38',  
              'f_39', 'f_40', 'f_41', 'f_42', 'f_43', 'f_44', 'f_45', 'f_46', 'f_47',  
              'f_48', 'f_49', 'target'],  
              dtype='object')
```

```
In [80]: correlation_metrix = df.corr()
```

```
In [81]: plt.figure(figsize = (20,16))  
sns.heatmap(correlation_metrix, annot = True, cmap = 'coolwarm', fmt = '.2f', linewidths= .5)  
plt.show()
```



```
In [82]: columns_of_interest = ['f_1', 'f_2', 'f_3', 'f_4', 'f_5', 'f_6', 'f_7', 'f_8', 'f_9', 'f_10',  
                                'f_11', 'f_12', 'f_13', 'f_14', 'f_15', 'f_16', 'f_17', 'f_18', 'f_19',  
                                'f_20', 'f_21', 'f_22', 'f_24', 'f_25', 'f_26', 'f_27', 'f_28', 'f_29',  
                                'f_30', 'f_31', 'f_32', 'f_33', 'f_34', 'f_35', 'f_36', 'f_37', 'f_38',  
                                'f_39', 'f_40', 'f_41', 'f_42', 'f_43', 'f_44', 'f_45', 'f_46', 'f_47',  
                                'f_48', 'f_49']  
  
fig, axes = plt.subplots(nrows=len(columns_of_interest), ncols=1, figsize=(10, 2 * len(columns_of_interest)))  
  
for i, column in enumerate(columns_of_interest):  
    df[column].plot(ax=axes[i], kind='hist', bins=30, legend=True)  
    axes[i].set_title(column)  
    axes[i].set_xlabel('')  
    axes[i].set_ylabel('')  
  
# Adjusting Layout  
plt.tight_layout()  
plt.show()
```



**04) Apply various Machine Learning techniques to predict the output in the target column, make use of Bagging and Ensemble as required, and find the best model by evaluating the model using Model evaluation techniques.**

#### Split the Data into Train and Test Sets

```
In [84]: X = df.drop(columns='target')  
y = df['target']
```

```
In [85]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```

### Model Selection and Training:

```
In [20]: rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
         rf_model.fit(X_train, y_train)
```

```
Out[20]: RandomForestClassifier
         RandomForestClassifier(random_state=42)
```

### Bagging and Ensemble Techniques:

```
In [21]: ada_model = AdaBoostClassifier(n_estimators=100, random_state=42)
         ada_model.fit(X_train, y_train)
```

```
Out[21]: AdaBoostClassifier
         AdaBoostClassifier(n_estimators=100, random_state=42)
```

### Model Evaluation:

```
In [22]: def evaluate_model(model, X_test, y_test):
         y_pred = model.predict(X_test)
         acc = accuracy_score(y_test, y_pred)
         return acc

         rf_accuracy = evaluate_model(rf_model, X_test, y_test)
         ada_accuracy = evaluate_model(ada_model, X_test, y_test)
```

## Q5) Save the best model and Load the model.

### Save and Load the Best Model:

```
In [23]: best_model = rf_model if rf_accuracy > ada_accuracy else ada_model
         joblib.dump(best_model, 'best_model.pkl')
```

```
Out[23]: ['best_model.pkl']
```

### Load the saved model:

```
In [24]: loaded_model = joblib.load('best_model.pkl')
```

## Q6) Take the original data set and make another dataset by randomly picking 20 data points from the oil spill dataset and applying the saved model to the same.

### Apply the Model to a Subset of Data:

```
In [25]: subset_df = df.sample(n=20, random_state=42)
         X_subset = subset_df.drop(columns=['target'])
         y_subset_true = subset_df['target']
         y_subset_pred = loaded_model.predict(X_subset)
```

### Display predictions on the subset data:

```
In [26]: print("Predictions on Subset Data:")
         print(y_subset_pred)

         Predictions on Subset Data:
         [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

## Analyze the characteristics of the subset data:

```
In [27]: print("Subset Data Info:")
print(subset_df.info())
```

```
Subset Data Info:
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20 entries, 321 to 244
Data columns (total 49 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   f_1      20 non-null      int64
 1   f_2      20 non-null      float64
 2   f_3      20 non-null      float64
 3   f_4      20 non-null      float64
 4   f_5      20 non-null      float64
 5   f_6      20 non-null      float64
 6   f_7      20 non-null      float64
 7   f_8      20 non-null      float64
 8   f_9      20 non-null      float64
 9   f_10     20 non-null      float64
10  f_11     20 non-null      float64
11  f_12     20 non-null      float64
12  f_13     20 non-null      float64
13  f_14     20 non-null      float64
```

## Compare feature distributions between training and subset data:

```
In [28]: print("\nTraining Data Describe:")
print(X_train.describe())
print("\nSubset Data Describe:")
print(X_subset.describe())
```

```
Training Data Describe:
      f_1      f_2      f_3      f_4      f_5      f_6 \
count  749.000000  749.000000  749.000000  749.000000  749.000000  749.000000
mean    82.544726   0.019476  -0.001653  -0.001427  -0.026036   0.022394
std    64.563112   1.103986   1.003222   1.006780   1.011358   1.110317
min     1.000000  -0.167197  -1.161999  -1.664992  -1.855452  -0.182650
25%    32.000000  -0.162018  -1.023017  -0.825101  -0.686443  -0.168367
50%    65.000000  -0.139749  -0.015288  -0.247899  -0.267365  -0.154474
75%   125.000000  -0.107640   0.875639   0.735508   0.659019  -0.119400
max   352.000000  16.601603   1.818446   3.547380   2.114766  18.423439

      f_7      f_8      f_9      f_10  ...      f_40 \
count  749.000000  749.000000  749.000000  749.000000  ...  749.000000
mean     0.003833  -0.030967   0.018737  -0.029306  ...   0.005495
std     1.008952   0.968020   1.097941   0.968225  ...   1.002029
min    -1.730915  -2.313347  -0.401040  -2.226754  ...  -1.557855
25%    -0.754642  -0.676864  -0.314797  -0.675806  ...  -0.762570
50%    -0.255099  -0.267047  -0.230515  -0.232678  ...  -0.401076
75%     0.728254   0.435498  -0.072058   0.432014  ...   0.611105
max     3.099313   4.224219  19.208377   5.749552  ...   1.840183

      f_41      f_42      f_43      f_44      f_45      f_46 \
count  749.000000  749.000000  749.000000  749.000000  749.000000  749.000000
mean     0.000000   0.000000   0.000000   0.000000   0.000000   0.000000
std     0.000000   0.000000   0.000000   0.000000   0.000000   0.000000
min     0.000000   0.000000   0.000000   0.000000   0.000000   0.000000
25%     0.000000   0.000000   0.000000   0.000000   0.000000   0.000000
50%     0.000000   0.000000   0.000000   0.000000   0.000000   0.000000
75%     0.000000   0.000000   0.000000   0.000000   0.000000   0.000000
max     0.000000   0.000000   0.000000   0.000000   0.000000   0.000000
```

## Analyze feature importance of the trained model:

```
In [29]: if hasattr(loaded_model, 'feature_importances_'):
feature_importance = loaded_model.feature_importances_
print("\nFeature Importance:")
for i, importance in enumerate(feature_importance):
    print(f"Feature {X_train.columns[i]}: {importance}")
else:
    print("The model does not support feature importance analysis.")
```

```
Feature Importance:
Feature f_1: 0.06281822414999656
Feature f_2: 0.019821470562562695
Feature f_3: 0.019501202791451514
Feature f_4: 0.03455306953609941
Feature f_5: 0.018332321495875964
Feature f_6: 0.02527086393365474
Feature f_7: 0.022456039846791957
Feature f_8: 0.022582236413578642
Feature f_9: 0.025628790818362734
Feature f_10: 0.0219858798777173
Feature f_11: 0.02595011178301022
Feature f_12: 0.016878204879060547
Feature f_13: 0.014379144251241332
Feature f_14: 0.016928888715298365
Feature f_15: 0.012609691320471583
Feature f_16: 0.023380394175093334
Feature f_17: 0.02392199288421556
Feature f_18: 0.017225081439483
Feature f_19: 0.0195702388954833
Feature f_20: 0.010105385431171218
Feature f_21: 0.01957605061545208
Feature f_22: 0.012141836290836728
Feature f_23: 0.02063551770100013
```